

# Socket Programming Basics

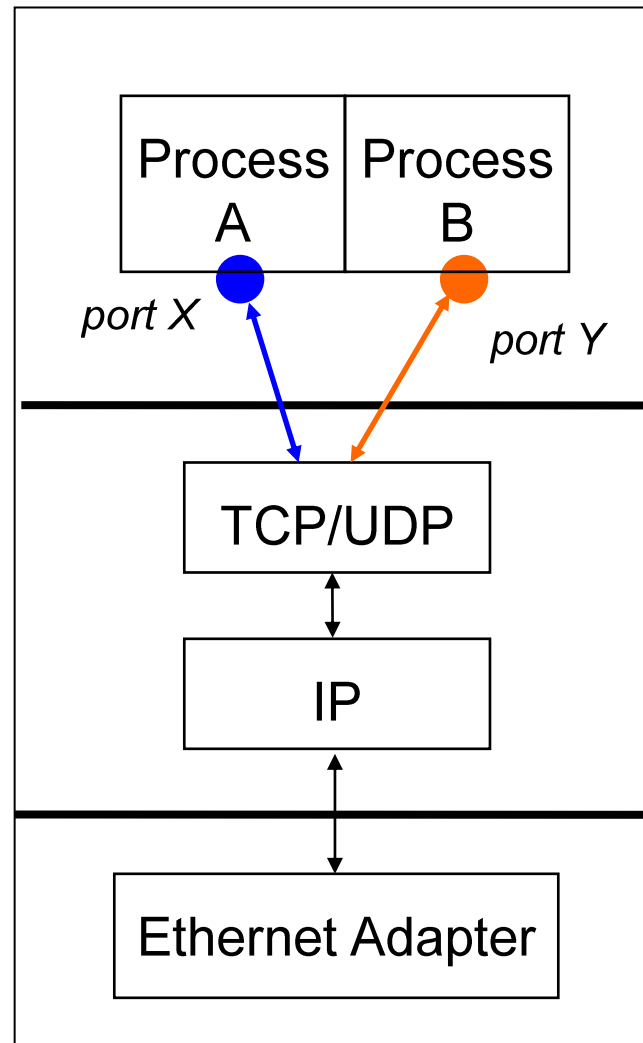
# Need of Network Sockets

- To support message communication between two Networking Applications running on same (may user process communication as well) or different hosts

# Properties of Network Socket

- Address of Computer / Node on Network?
  - **IP Address** (IPv4)
    - A 32 bit **unique** address on your network (Ex. 8.8.8.8)
    - Handled by L3 (IP/Internet Layer) of 5/7 layer protocol stack
- Address of Network Application on a Computer
  - **Port Number** [Logical]
    - A 16 bit unique identifier of Network Application
    - Handled by L4 (Transport Layer) of a 5/7 layer protocol stack
- Type of Communication
  - **TCP** (Connection Oriented / Reliable) / **UDP** (Connectionless / Best Effort)

# Socket Identification




# Socket Identification

- Can a machine has more than one IP address ?
  - YES, Each Network Interface Card (NIC) has a unique IP addresses (Ethernet and Wi-Fi NIC are assigned different IP address and both can be used at same time)
  - One can create multiple virtual Network Interfaces with each assigned different IP address associated with same NIC.
- Can one Socket (server) be made to associated with more than one IP
  - Either One or ALL
  - What if we want multiple ?
- Can two sockets be made to listen on same port
  - No, if the bind IP and Transport Layer protocol are same
  - Yes, if any of above is different

# Client-Server Communication

- Client "sometimes on"
  - Initiates a request to the server when interested
  - E.g., Web browser on your laptop or cell phone
  - Doesn't communicate directly with other clients
  - Needs to know server's address
- Server is "always on"
  - Handles services requests from many client hosts
  - E.g., Web server for the [www.iitk.ac.in](http://www.iitk.ac.in) Web site
  - Doesn't initiate contact with the clients
  - Needs fixed, known address



# Client-Server Communication

## Stream Sockets (TCP): Connection-oriented

### Server

socket()

Create a socket

bind()

Bind the socket  
(port/IP/protocol)

listen()

Listen for client  
(Wait for incoming connections)

accept()

Accept connection  
(Blocking Method)

recv()

Receive Request / Data

send()

Send response / Data

### Client

Create a socket

socket()

Connect to server

connect()

Send the request  
/ data

send()

Receive response /  
data

recv() <sub>7</sub>

establish connection  
(Three way handshake)

data (request)

data (reply)

# Implementation (Server)

- Create a Socket [Creating the Endpoint for Socket]
  - ***int sockid = socket(family, type, protocol);***
  - *Defined in sys/socket.h*
  - ***sockid:** socket descriptor, an integer (like a file-handle)*
  - ***family:** integer, communication domain, e.g.,*
    - *AF\_INET, IPv4 protocols, Internet addresses (typically used)*
    - *AF\_UNIX, Local communication, File addresses*
  - ***type:** communication type*
    - *SOCK\_STREAM - reliable, 2-way, connection-based service [TCP]*
    - *SOCK\_DGRAM - unreliable, connectionless, messages of maximum length [UDP]*
  - ***protocol:** specifies protocol*
    - *IPPROTO\_TCP IPPROTO\_UDP*
    - *usually set to 0 (i.e., use default protocol)*
  - *upon failure returns -1*
  - *Ex: **int sockid = socket(AF\_INET, SOCK\_STREAM, 0);***



# Specifying Address

- Socket API defines a generic data type for addresses:

```
struct sockaddr {  
    unsigned short sa_family;    /* Address family (e.g. AF_INET) */  
    char sa_data[14];           /* Family-specific address information */  
}
```

- Particular form of the sockaddr used for TCP/IP addresses:

```
struct in_addr {  
    unsigned long s_addr        /* Internet address (32 bits) */  
}  
  
struct sockaddr_in {  
    unsigned short sin_family;  /* Internet protocol (AF_INET) */  
    unsigned short sin_port;    /* Address port (16 bits) */  
    struct in_addr sin_addr;    /* Internet address (32 bits) */  
    char sin_zero[8];          /* Not used */  
}
```

- `sockaddr_in` can be casted to a `sockaddr`

# Assigning address to Socket

```
int status = bind(sockid, &addr, size);
```

- **sockid**: integer, socket descriptor
- **addr**: struct sockaddr, the (IP) address and port of the machine
  - for TCP/IP server, internet address is usually set to INADDR\_ANY, i.e., chooses ALL incoming interface
  - or can be specified a particular IP
- **size**: the size (in bytes) of the addr structure
- **status**: 0 if successful bind, -1 otherwise
- Ex:

```
struct sockaddr_in addr;  
addr.sin_family = AF_INET;  
addr.sin_port = htons(5100);  
addr.sin_addr.s_addr = htonl(INADDR_ANY);  
bind(sockid, (struct sockaddr *) &addr, sizeof(addr10))
```

# Start Listening for Connections

```
int status = listen(sockid, queueLimit);
```

- **Instructs TCP protocol implementation to listen for connections**
  - sockid: integer, socket descriptor
  - queueLen: integer, # of active participants that can “wait” for a connection while server is busy serving previously arrived client.
  - status: 0 if listening, -1 if error

# Accepting connection request

```
int newsockid = accept(sockid, &clientAddr, &addrLen);
```

- newsockid: integer, the new socket (used for data-transfer)
- sockid: integer, the orig. socket (being listened on)
- clientAddr: struct sockaddr, address of the active participant
  - filled in upon return with the details of client information (which is available in the IP packet received by server)
  - addrLen: sizeof(clientAddr): value/result parameter
- **accept** method call is a blocking call. Program thread will keep waiting till it receives a connection request from client.

# Data Communication

- Receiving Data

**n = read(newsockid, buffer, count);**

- newsockid : integer value returned by accept function call
- buffer: buffer to store the read value
- count: Number of bytes to be read
- n - Number of blocks actually read (may be less than count)

- Sending Data

**n = write(newsockid, buffer, count);**

- newsockid : integer value returned by accept function call
- buffer: data to be written
- count: Number of bytes to be read
- n - Number of blocks actually read (may be less than count)

- **fdopen** : can be used to make socket stream behave like FILE stream so that fread / fwrite and similar functions can be used.

# Implementation (Client)

- Create the socket (end point for communication) same as it was created for server.

**int sockid = socket(family, type, protocol);**

- Client executes a **connect** method call to send socket creation request to the server.

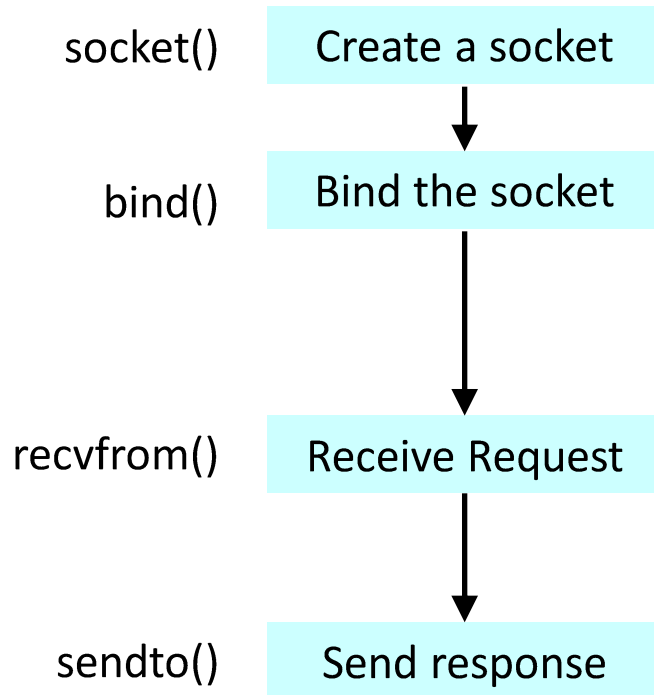
**int status = connect(sockid, &serverAddr, addrlen);**

- sockid: integer, socket to be used in connection
  - serverAddr: struct sockaddr: address of the passive participant
  - addrlen: integer, sizeof(name)
  - status: 0 if successful connect, -1 otherwise
- On Successful connection sockid can be used to read and write data in the same way it was implemented for Server

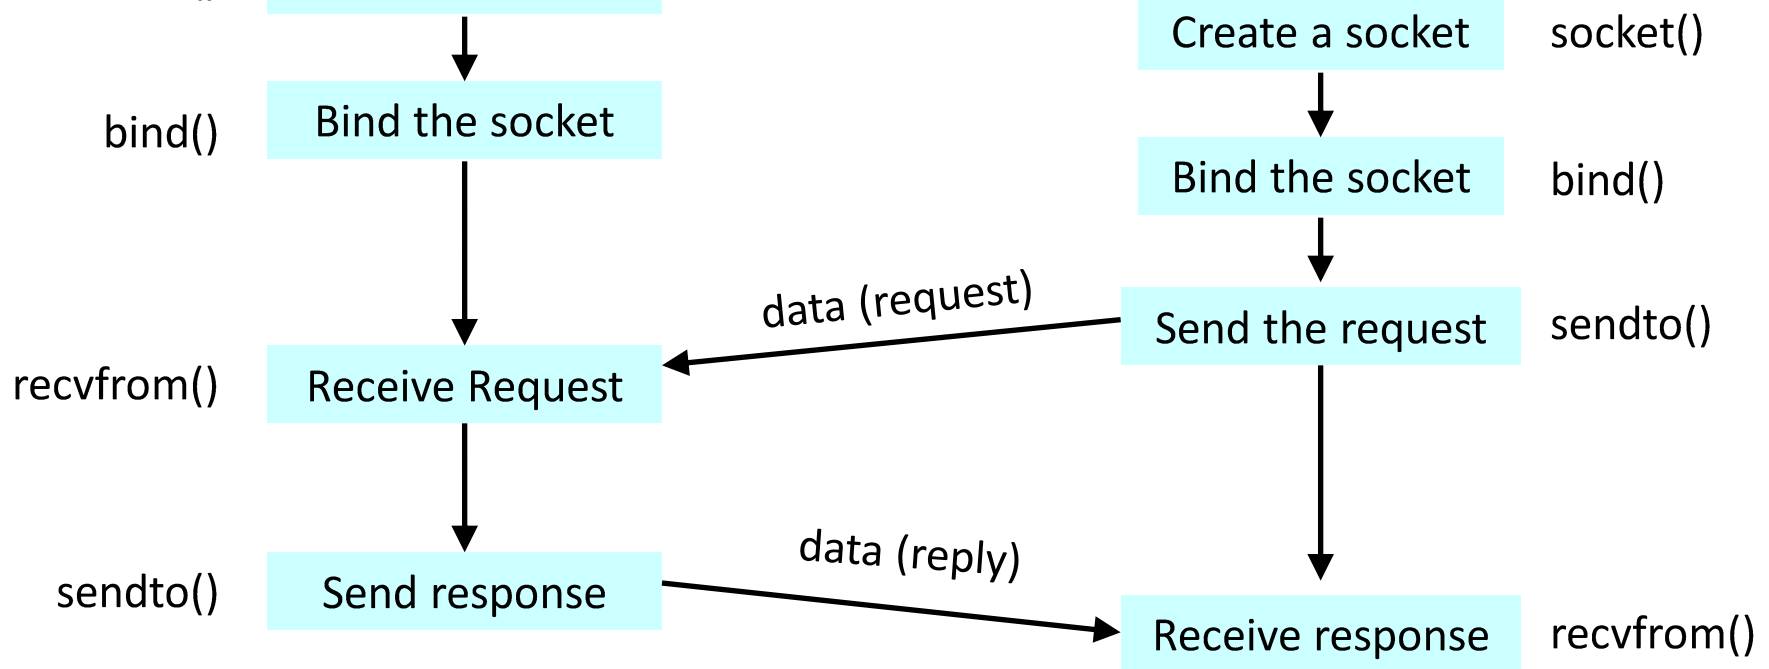
# Client-Server Communication

## Datagram Sockets (UDP): Connectionless

### Server



### Client



# Blocking vs. Non Blocking Socket

- Blocking Socket:
  - Clients will be served sequentially
- Non Blocking Socket:
  - Clients may be served in parallel
- Serving Clients in parallel with Blocking Socket
  - Create a child process or thread each time accept method is called. Child process will serve the client while parent will return to listening mode.



# Concurrent Server - Example

```
while(1){  
  connfd = accept(listenfd, ...); /* blocking call */  
  if ( (pid = fork()) == 0 ) {  
    close(listenfd); /* child closes listening socket */  
    /**process the request doing something using connfd ***/  
    close(connfd);  
    exit(0); /* child terminates */  
  }  
  close(connfd); /*parent closes connected socket*/  
}
```

- Questions ?